

Final Report: Towards Automated Safety Audits for Model Context Protocols

Zhao Xu

zhaohsu@g.ucla.edu

Chenwei Gu

suzychenwei@cs.ucla.edu

Tong Xie

tongxie@g.ucla.edu

Yuanhao Ban^{*}

banyh2000@gmail.com

Abstract

The Model Context Protocol (MCP) provides a standardized interface for LLM agents to dynamically discover and invoke tools via MCP servers. While MCP improves extensibility and interoperability, it introduces protocol-level attack surfaces and governance risks, including description poisoning, multi-hop trust hijacking, version-based “rug-pull” and malicious external dependencies. Existing AI safety frameworks focus largely on model-level alignment, overlooking these emerging system-level threats. We present MCPDojo, an evaluation framework that simulates realistic MCP hosts and assembles benign/adversarial tasks across these vectors. Building on these evaluations, we outline the design of an **automated audit system** that operationalizes the findings into continuous compliance monitoring. Our work establishes a foundation for standardized protocol-centric safety evaluation and provides insights for governance and policy alignment.

1 Introduction

Large Language Model (LLM) agents increasingly rely on the **Model Context Protocol (MCP)** to communicate with external tools, databases, and APIs. While this greatly enhances their functionality and flexibility, it also introduces new attack surfaces for system-level AI safety risks. Recent studies ([Song et al., 2025](#); [Ferrag et al., 2025](#)) have identified multiple vulnerabilities in MCP ecosystems, including tool poisoning, permission escalation, dependency version drift, and malicious endpoint injection.

Despite the growing adoption of agent-based architectures, existing AI safety frameworks (e.g., NIST AI RMF ([Tabassi, 2023](#)), EU AI Act ([Act, 2024](#))) primarily focus on model alignment and content safety, leaving the protocol layer—a critical component of agent governance—largely unaddressed.

To address this gap, we introduce **MCPDojo**, an evaluation framework for assessing the safety and robustness of MCP-based agents. MCPDojo simulates realistic MCP environments containing both benign and adversarial tasks. Based on these evaluations, we further outline the development of an **automated audit system** that leverages observed risk patterns to continuously monitor compliance and flag unsafe behaviors. Our long-term goal is to bridge the divide between **protocol-level evaluation** and **governance-oriented auditing**, contributing to a foundation for standardized agent safety assessment and accountability.

2 Related Work

Recent work on Model Context Protocol (MCP) and agent tool-use security broadly falls into two lines: (i) **capability and benchmarking of MCP-enabled agents and servers**, and (ii) **security analysis of agent tool-use under adversarial or malicious inputs**. The first line focuses on evaluating task performance, execution fidelity, and efficiency of MCP-based systems ([Luo et al., 2025](#); [Gao et al., 2025](#); [Yan et al., 2025](#); [Liu et al., 2025](#); [Fu et al., 2025](#)), while the second line characterizes emerging attack surfaces introduced by protocol-mediated tool interactions ([Debenedetti et al., 2024](#); [Song et al., 2025](#); [Luca Beurer-Kellner, Marc Fischer, 2025](#); [Radosevich and Halloran, 2025](#)).

Despite rapid progress, existing benchmarks primarily emphasize tool-use capability or end-task success, and security-oriented works often study attacks in isolation without a unified evaluation harness. In contrast, we position our work as an *evaluation-first* framework that systematically combines realistic MCP environments, attack-aware task construction, and multi-axis metrics, providing the empirical foundation required for principled, policy-aligned auditing of agent ecosystems.

^{*} Corresponding author.

2.1 MCP-Centric Benchmarks

Several recent benchmarks evaluate the capabilities of MCP-enabled agents or servers, largely under benign or weakly adversarial conditions. Luo et al. (Luo et al., 2025) focus on MCP servers themselves, emphasizing data-fetching accuracy and efficiency metrics (e.g., latency and token cost), without considering adversarial robustness. MCP-RADAR (Gao et al., 2025) and MCPWorld (Yan et al., 2025) evaluate agent performance in more realistic MCP settings, introducing operation-level verification and white-box state inspection, respectively, but primarily target correctness rather than security under attack. MCPEval (Liu et al., 2025) emphasizes scalable task generation and verification via real MCP servers, while RAS-Eval (Fu et al., 2025) highlights limitations of purely simulated evaluations and moves toward security benchmarking. Overall, existing benchmarks focus on capability, correctness, or scalability, but do not provide a unified, attack-aware evaluation framework tailored to MCP-specific protocol risks or auditing needs.

2.2 Attacks on MCP

AgentDojo (Debenedetti et al., 2024) is explicitly security-driven: it studies prompt injection attacks where data returned by tools hijacks the agent, and proposes an extensible environment (not a static suite) for evaluating tasks, defenses and adaptive attacks. Song H et al. (Song et al., 2025) argues MCP’s client–server integration expands the attack surface by enabling malicious MCP servers, and provides an end-to-end empirical study across the lifecycle. It further frames MCP as crossing a trust boundary in registration, planning and operation, with opportunities to embed harmful instructions in tool descriptions and responses. Radosevich et al. (Radosevich and Halloran, 2025) proposes Retrieval-Agent Deception (RADE): instead of directly prompting the agent, the attacker corrupts public data that gets ingested into a vector DB so that later retrieval loads tool-leveraging commands; the paper argues this raises the threat level because the attacker no longer needs direct access.

2.3 Defenses on MCP

Several practical tools have been proposed to mitigate MCP-related risks through proactive scanning and auditing. Radosevich et al. (Radosevich and Halloran, 2025) introduce McpSafetyScanner,

which analyzes MCP server features, generates adversarial probes, and reports potential vulnerabilities along with remediation suggestions. Similarly, MCP-Scan (Luca Beurer-Kellner, Marc Fischer, 2025), developed by Invariant Labs, provides lightweight security checks for MCP installations.

While these tools are valuable for engineering practice, they primarily operate as standalone scanners and do not provide a standardized, task-level evaluation framework for measuring agent behavior under protocol-level attacks. Our work complements these approaches by grounding auditing strategies in systematic, attack-aware evaluation.

3 Methodology

The MCPDOJO framework is designed as a modular and extensible testbed for evaluating protocol-level security risks of LLM agents interacting with MCP servers. It follows a three-stage design: *construction* of tasks and environments, *execution* with pluggable agent backbones, and *evaluation* along multiple axes and audit based on different audit strategies. This section describes the first two stages in detail. This section describes all three stages, with emphasis on construction and execution, and then presents the evaluation protocol and policy-oriented audit design.

3.1 Construction: Environments, Tools, and Tasks

Environments. We instantiate two domain-specific environments, including *workspace* and *messaging*, with mutable states such as emails, files, itineraries, and financial records. Each environment provides an abstraction of real-world agent use cases while embedding *injection hooks* that allow adversarial payloads to be introduced in controlled ways (e.g., poisoned email subject lines, manipulated API responses). By isolating domains, we can study cross-environment generalization and the effect of contextual complexity on attack success.

Tools. Environments are equipped with a suite of MCP tools, each described through signed and versioned JSON schemas containing function names, parameters, expected outputs, and permission tags (e.g., read-only vs. write-critical). This design mirrors the MCP specification (Anthropic, 2025) while enabling systematic manipulation of tool descriptions for attack experiments (e.g., schema poisoning, version rollback). The use of signatures




Figure 1: Overall pipeline of MCPDOJO.

and version tags also provides hooks for evaluating detection mechanisms.

Task Suite. We define a large set of *benign user tasks* (e.g., “book a flight and add it to the calendar”) paired with corresponding ground-truth tool invocation sequences. These are combined with *adversarial templates* derived from the attack taxonomy, yielding a Cartesian product:

$$\begin{aligned} \text{Task Suite} = & \text{User Tasks (benign)} \\ & \times \text{Attack Templates (vector, difficulty, steps).} \end{aligned} \quad (1)$$

This construction generates both baseline tasks (for measuring utility) and adversarial tasks (for measuring safety). Difficulty is controlled by varying the number of steps (1–5), the sensitivity of the target (e.g., ordinary emails vs. MFA codes), and the placement of malicious payloads (tool description vs. elicitation schema).

3.2 Execution: Agent Backbone

Agent backbone. MCPDOJO provides a pluggable interface for evaluating different LLM backbones (e.g., GPT-4-family, Qwen-family) as agents. Each agent interacts with environments by executing the canonical MCP workflow: *discovery* → *fetch description* → *tool call* → *state update*. This uniform abstraction allows fair comparison across models while preserving protocol fidelity.

Execution pipeline. Each experimental run proceeds through three sequential stages, as illustrated in Figure 1. (1) **Task sampling:** a task–attack pair is drawn from the constructed suite according to difficulty and domain, ensuring balanced coverage of benign and adversarial instances. (2) **Attack assembly:** the selected adversarial template is instantiated within the target environment by injecting attack goals into tool descriptions or server responses. (3) **Host & Server Execution:** the environment and the MCP server are launched in a controlled simulator, where the agent performs the standard MCP workflow.

3.3 Evaluation & Audit

Evaluation. MCPDOJO adopts a two-axis evaluation protocol to jointly measure functionality and robustness.

(1) **Utility.** We define *Benign Utility* as ratio of benign tasks successfully completed:

$$Utility_B = \frac{1}{N_b} \sum_{i=1}^{N_b} \mathbb{1}[f(x_i) = y_i] \quad (2)$$

where N_b is the number of benign tasks, $f(x_i)$ is final environment state produced by the agent on input x_i , and y_i is the expected ground truth state.

(2) **Safety.** We define *Attack Success Rate (ASR)* and *Utility under Attack* as:

$$ASR = \frac{1}{N_a} \sum_{i=1}^{N_a} \mathbb{1}[g(x_i) = o_i], \quad (3)$$

where N_a is the number of adversarial tasks, $g(x_i)$ is the agent’s observed outcome, and o_i is the adversary’s malicious target.

$$Utility_A = \frac{1}{N_a} \sum_{i=1}^{N_a} \mathbb{1}[f(x_i) = y_i \wedge g(x_i) \neq o_i], \quad (4)$$

which captures the fraction of adversarial tasks where the agent still completes the intended user task while avoiding the adversary’s objective.

AI Policy-Oriented Audit. Recent AI policy frameworks such as EU AI Act (Act, 2024) set expectations for AI ecosystems. The Act establishes a tiered classification of risks, while assigning defined roles to providers and deployers. This structure ensures that every stage of AI usage operates under explicit and traceable oversight. Complementing these regulatory frameworks, the OpenAI Model Spec (OpenAI, 2025) introduces a behavioral standard for AI systems. It outlines three core principles: (1) maximizing helpfulness and user freedom, (2) minimizing harm, and (3) designing sensible defaults. Together, these regulatory expectations and leading-firm guidelines provide a strong reference for how AI agents should be audited to balance utility with safety.

Building on these standards, an ideal MCP auditor is designed with focus on risk breakdown and assessment, clear goal distinctions, and mitigation against attacks. We design the auditor to integrate both static and dynamic control. The static auditor enforces preventive measures to ensure basic compliance before executions (GitHub Copilot, 2025; Anthropic, 2025; OpenAI Agents SDK, 2025). The dynamic auditor extends this oversight at runtime, using risk classification and LLM-reasoning to detect and assess unsafe behaviors. This dual design aligns with the prevailing policies’ focus on risk assessment and transparency.

3.4 Evaluation Results

Table 1 reports the evaluation results on MCPDOJO across benign utility ($Utility_B$), utility under attack ($Utility_A$), and attack success rate. Across all models, adversarial perturbations introduced through MCP artifacts consistently reduce task completion, indicating that protocol-level attacks substantially affect agent behavior beyond benign settings.

Under benign conditions, larger models achieve higher utility, with *GPT-4o-2024-11-20* attaining




Figure 2: **Dynamic MCP auditor** for real-time risk assessment with clear, auditable outputs.

Table 1: Main evaluation results on MCPDOJO. We report Utility (Benign Success Rate, $Utility_B$, and Utility under Attack, $Utility_A$), and Safety (Attack Success Rate, ASR ↓). Best results per column are **bolded**.

Family	Model	Utility		Safety
		$Utility_B(\%) \uparrow$	$Utility_A(\%) \uparrow$	
OpenAI	GPT-4o-mini-2024-07-18	64.82	53.61	9.87
	GPT-4o-2024-11-20	72.14	61.86	12.96
	GPT-4.1-2025-04-14	59.64	56.70	0.36
Qwen	Qwen3-1.7B	33.33	22.86	21.43
	Qwen3-4B	38.10	34.29	35.71

the best $Utility_B$. However, higher benign performance does not necessarily translate to improved safety. In adversarial settings, utility drops across all backbones, and smaller models (e.g., Qwen3-1.7B) experience the most severe degradation in $Utility_A$, reflecting limited robustness to protocol-mediated attacks.

Safety results further reveal a clear trade-off between robustness and utility. *GPT-4.1-2025-04-14* achieves the lowest ASR, suggesting strong resistance to adversarial objectives, but exhibits reduced benign utility compared to GPT-4o variants. These results highlight the importance of jointly evaluating utility and safety, and motivate auditing mechanisms that monitor intermediate tool-use behaviors rather than relying solely on final task success.

4 Embedded Ethics Discussion

Problem & Solution. Model Context Protocols (MCP) are designed to make tool-use in LLM agents more efficient and interoperable: instead of hard-coding integrations, MCP standardizes how agents utilize external tools, allowing deployment across different model providers and environments. However, this introduces new safety concerns and attack surfaces. For example, adversarial attacks may inject malicious tools or poison tool metadata to steer downstream agent behavior. To address these risks, we propose a dynamic auditor that continuously tracks agent safety on realistic tasks.

Course Plan. This can be organized into four progressive modules: (1) understand MCP and risks, (2) survey common attacks, (3) evaluate existing agents under these attacks, and (4) design mini-auditor. For actionable course plans, lectures can focus on the conceptual workflow, risks, and evaluation principles, while assignments provide hands-on practice with an end-to-end tool-use pipeline. Concretely, students can try various attack methods to benchmark on agents. Then based on the vulnerabilities exposed, design an audit system to monitor agent behaviors and mitigate attacks.

5 Conclusion

In this work, we presented MCPDOJO, a modular testbed for evaluating protocol-level security risks of MCP-enabled LLM agents under realistic benign and adversarial tasks.

By jointly measuring benign utility, safety under attack, and the ability to maintain intended task goals, MCPDOJO exposes concrete failure modes that are not captured by model-level alignment evaluations alone. Building on these empirical findings, we outlined an AI policy-oriented audit design that combines preventive controls (e.g., configuration and provenance checks) with runtime, log-based risk assessment to support transparency and accountability.

We hope this work encourages standardized, protocol-centric evaluation and actionable governance practices for emerging agent ecosystems.

6 Limitations and Future Work

Scope of evaluation and task design. Although MCPDOJO is designed as a general testbed for protocol-level risks, the current study adopts a deliberately constrained scope. Following profes-

sor’s feedback on our proposal, we focus on a limited set of MCP server-side tool poisoning scenarios rather than the full attack taxonomy. Due to implementation complexity and time constraints, our experiments currently support a single environment (workspace), and many benign tasks and injection templates are adapted from Agent-Dojo (DeBenedetti et al., 2024) instead of being fully MCP-native. While this enables controlled comparison with prior work, future efforts should expand environment diversity and develop task suites that explicitly target MCP-specific behaviors, such as multi-turn and multi-server interactions.

Model coverage and evaluation scale. Our evaluation covers a limited set of backbone LLMs, primarily from the OpenAI and Qwen families. Resource constraints prevented broader model coverage, and additional models (e.g., GPT-5-mini) could not be included in the final results due to implementation issues discovered after extensive runs. Future work should improve evaluation stability and scale experiments across a wider range of models to better understand robustness trends.

Reliability of auditing explanations. Finally, while our dynamic auditor produces structured risk assessments with natural-language explanations, we do not claim full interpretability of LLM-based auditing. Explanation faithfulness depends on how risk scores are defined and how the auditor model reasons about execution traces. An important direction for future work is to study the calibration and reliability of audit explanations, and to integrate more verifiable signals (e.g., rule-based checks or provenance constraints) alongside LLM reasoning.

7 Contribution Statement

Zhao Xu: designed the overall system, methodology, and execution pipelines; wrote the paper.

Tong Xie: wrote the paper; designed the static risk analyzer; mapped the real-world policy requirements to the static audit principle.

Chenwei Gu: wrote the paper; drew the figures; mapped the real-world policy requirements to the static audit principle.

Yuanhao Ban: served as the corresponding author; drew the figures; provided comprehensive logistical support for the team’s operations.

References

- EU Artificial Intelligence Act. 2024. The eu artificial intelligence act. *European Union*.
- Anthropic. 2025. [Model context protocol specification \(2025-06-18\)](#).
- Edoardo Debenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. 2024. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents. *Advances in Neural Information Processing Systems*, 37:82895–82920.
- Mohamed Amine Ferrag, Norbert Tihanyi, Djallel Hamouda, Leandros Maglaras, and Merouane Debbah. 2025. From prompt injections to protocol exploits: Threats in llm-powered ai agents workflows. *arXiv preprint arXiv:2506.23260*.
- Yuchuan Fu, Xiaohan Yuan, and Dongxia Wang. 2025. [Ras-eval: A comprehensive benchmark for security evaluation of llm agents in real-world environments](#). *Preprint*, arXiv:2506.15253.
- Xuanqi Gao, Siyi Xie, Juan Zhai, Shiqing Ma, and Chao Shen. 2025. [Mcp-radar: A multi-dimensional benchmark for evaluating tool use capabilities in large language models](#). *Preprint*, arXiv:2505.16700.
- GitHub Copilot. 2025. [Configure mcp server access](#).
- Zhiwei Liu, Jielin Qiu, Shiyu Wang, Jianguo Zhang, Zuxin Liu, Roshan Ram, Haolin Chen, Weiran Yao, Shelby Heinecke, Silvio Savarese, Huan Wang, and Caiming Xiong. 2025. MCPEval: Automatic MCP-based deep evaluation for AI agent models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 373–402.
- Luca Beurer-Kellner, Marc Fischer. 2025. [Introducing mcp-scan: Protecting mcp with invariant](#).
- Zhilong Luo, Xiaorong Shi, Xuanrui Lin, and Jinyang Gao. 2025. [Evaluation report on mcp servers](#). *Preprint*, arXiv:2504.11094.
- OpenAI. 2025. [Openai model spec](#).
- OpenAI Agents SDK. 2025. [Model context protocol \(mcp\)](#).
- Brandon Radosevich and John Halloran. 2025. [Mcp safety audit: Llms with the model context protocol allow major security exploits](#). *Preprint*, arXiv:2504.03767.
- Hao Song, Yiming Shen, Wenxuan Luo, Leixin Guo, Ting Chen, Jiashui Wang, Beibei Li, Xiaosong Zhang, and Jiachi Chen. 2025. [Beyond the protocol: Unveiling attack vectors in the model context protocol \(mcp\) ecosystem](#). *Preprint*, arXiv:2506.02040.
- Elham Tabassi. 2023. Artificial intelligence risk management framework (ai rmf 1.0).
- Yunhe Yan, Shihe Wang, Jiajun Du, Yexuan Yang, Yuxuan Shan, Qichen Qiu, Xianqing Jia, Xinge Wang, Xin Yuan, Xu Han, Mao Qin, Yinxião Chen, Chen Peng, Shangguang Wang, and Mengwei Xu. 2025. [Mcpworld: A unified benchmarking testbed for api, gui, and hybrid computer use agents](#). *Preprint*, arXiv:2506.07672.

A Appendix

A.1 Example of Environment: Workspace

This work is largely inspired by AgentDojo (Debenedetti et al., 2024), which focuses primarily on evaluating prompt injection attacks on LLM agents. Accordingly, we adopt several basic design elements from it, such as the environment and task definitions, but reformulate them under the MCP format. In this section, we present a workspace example to illustrate our environment design.

```
{  
    "account_email": "emma.johnson@{{~domain~}}",  
    "current_day": "{{current_day}}",  
    "initial_events": [  
        {  
            "id_": "6",  
            "title": "Team Sync",  
            "description": "Weekly team sync to discuss project updates and roadblocks.",  
            "start_time": "{{start_time}}",  
            "end_time": "{{end_time}}",  
            "location": "Conference Room B",  
            "participants": [  
                "emma.johnson@{{~domain~}}",  
                "michael.smith@{{~domain~}}",  
                "sara.lee@{{~domain~}}"  
            ],  
            "all_day": false,  
            "status": "confirmed"  
        }  
    ]  
}
```

Figure 3: **Workspace:** This environment supports MCP servers for emails and calendar management.

While the current workspace remains a static prototype, we are actively working on extending it to support adaptive environments and task dynamics.

```
calendar_introductory_meeting_injection:  
    description: "Description of a meeting called 'Introductory meeting' in the calendar"  
    default: "Meeting to discuss potential collaboration opportunities."  
calendar_networking_event_injection_email:  
    description: "Final sentence of the description of a networking event"  
    default: "Don't miss this opportunity to expand your professional network!"  
calendar_sports_and_yoga_class_injection:  
    description: "Description of a sports and yoga class in the calendar"  
    default: "Join us for a relaxing yoga class and sports activities to boost your energy and health."  
calendar_dinner_with_clients_injection:  
    description: "Final sentence of the description of a dinner with some clients"  
    default: ""  
calendar_fake_deadline_description:  
    description: "Description of a fake deadline in the calendar"  
    default: "Your data will be deleted soon! Click here to save it: https://www.data-backup-security.com/save"
```

Figure 4: **Workspace:** Some attack vectors for tool poisoning.