

Saadia Gabriel

Email: skgabrie@cs.ucla.edu

Website: <https://saadiagabriel.com/>

PI: UCLA Misinformation, AI and Responsible Society Lab

Employment History

- 2024 – ····  **University of California, Los Angeles**
Samueli School of Engineering, Computer Science
Assistant Professor
- 2023-2024  **New York University**
Center for Data Science
Faculty Fellow/Assistant Professor
- 2023  **Massachusetts Institute of Technology**
Computer Science & Artificial Intelligence Laboratory
Postdoctoral Fellow with Marzyeh Ghassemi
- 2020, 2021  **Microsoft Research**
AI/NLP Research Intern
- 2019-2020  **Allen Institute for Artificial Intelligence**
Mosaic Team Research Intern
- 2019  **SRI International**
Computer Vision & Learning Research Intern
- 2015, 2016  **University of Massachusetts, Amherst**
Data Science Research Assistant

Education

- 2017 – 2023  **University of Washington**, PhD in Computer Science & Engineering.
Advisors: Yejin Choi and Franziska Roesner
- 2013 – 2017  **Mount Holyoke College**, BA (*summa cum laude*) in Computer Science & Mathematics.
Thesis Advisor: Daniel Sheldon

Awards & Honors

PI Fellowships & Awards

-  Google Research Scholar, 2025
-  Forbes 30 Under 30, Science 2024 List
-  Outstanding Reviewer, ACL 2020 and NAACL 2022
-  Google-Leap Fellowship, 2021
-  Phi Beta Kappa, 2017
-  Weaver Award for Computer Science and Math, Mount Holyoke College, 2017
-  David Notkin Endowed Graduate Fellowship, University of Washington, 2017
-  ARCS Foundation Fellowship, 2017

Paper Awards

-  SAC Highlights Award, *African Languages Lab*, ACL 2026
-  Editor's Pick, *Mental-LLM*, ACM IMWUT 2024

Awards & Honors (continued)

- 📌 MIT Generative AI Impact Award, *Generative AI in the Era of "Alternative Facts"*, 2023
- 📌 Best Paper, *Social Bias Frames*, West Coast NLP Summit 2020
- 📌 Best Paper Nomination, *Early Fusion for Goal Directed Robotic Vision*, IROS 2019
- 📌 Best Paper Nomination, *The Risk of Racial Bias in Hate Speech Detection*, ACL 2019

Selected Recent Research Publications

Preprints and Working Papers

- 1 S. Afroogh *et al.*, “Beyond explainable AI (XAI): An overdue paradigm shift and post-XAI research directions,” 2026. 🔗 URL: <https://arxiv.org/abs/2602.24176>.
- 2 M. Beikzadeh, Y. Salmanpour, A. Suvarna, *et al.*, “Multi-objective alignment of language models for personalized psychotherapy,” 2026. 🔗 URL: <https://api.semanticscholar.org/CorpusID:285726375>.
- 3 C. Chance, R. Pattichis, A. Subramonian, *et al.*, “TYKYK (but AI doesn’t): Automated content moderation does not capture communities’ heterogeneous attitudes towards reclaimed language,” 2026. arXiv: 2604.16654 [cs.CL]. 🔗 URL: <https://arxiv.org/abs/2604.16654>.
- 4 S. M. Issaka, E. R. Gonzalez, L. Liu, *et al.*, “Translation as a scalable proxy for multilingual evaluation,” 2026. 🔗 URL: <https://api.semanticscholar.org/CorpusID:284909942>.
- 5 S. Rahman, J. Shen, A. Mordvina, H. Palangi, S. Gabriel, and P. Izmailov, “When can LLMs learn to reason with weak supervision?” 2026. 🔗 URL: <https://arxiv.org/abs/2604.18574>.
- 6 A. Suvarna, K. Phan, M. Beikzadeh, H. Bansal, and S. Gabriel, “Supernova: Eliciting general reasoning in llms with reinforcement learning on natural instructions,” 2026. 🔗 URL: <https://arxiv.org/abs/2604.08477>.
- 7 S. Rahman, L. Y. Jiang, S. Gabriel, Y. Aphinyanagphongs, E. K. Oermann, and R. Chunara, “Generalization in healthcare AI: Evaluation of a clinical large language model,” 2024. 🔗 URL: <https://api.semanticscholar.org/CorpusID:267750868>.

Journal Articles

- 1 L. Jiang, J. D. Hwang, C. Bhagavatula, *et al.*, “Investigating machine moral judgement through the Delphi experiment,” *Nature Machine Intelligence*, 2025. 🔗 URL: <https://arxiv.org/abs/2110.07574>.
- 2 X. Xu, B. Yao, Y. Dong, *et al.*, “Mental-LLM: Leveraging Large Language Models for Mental Health Prediction via Online Text Data,” *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, vol. 8, no. 1, Mar. 2024. 🔗 DOI: 10.1145/3643540.

Conference and Workshop Proceedings

- 1 E. K. Guha, R. Marten, S. S. Keh, *et al.*, “OpenThoughts: Data recipes for reasoning models,” in *ICLR*, Apr. 2026. 🔗 URL: <https://api.semanticscholar.org/CorpusID:279154475>.
- 2 S. Issaka *et al.*, “The African Languages Lab: A collaborative approach to advancing low-resource African NLP,” in *ACL*, Jul. 2026. 🔗 URL: <https://arxiv.org/abs/2410.09949>.
- 3 G. Liu, V. Le, S. Rahman, E. Kreiss, M. Ghassemi, and S. Gabriel, “MOSAIC: Modeling social AI for content dissemination and regulation in multi-agent simulations,” in *EMNLP*, 2025. 🔗 URL: <https://arxiv.org/abs/2504.07830>.
- 4 S. Rahman, S. M. Issaka, A. Suvarna, *et al.*, “AI Debate Aids Assessment of Controversial Claims,” in *NeurIPS*, 2025. 🔗 URL: <https://api.semanticscholar.org/CorpusID:279119478>.

- 5 S. Rahman, L. Jiang, J. Shiffer, *et al.*, “X-teaming: Multi-turn jailbreaks and defenses with adaptive multi-agents,” in *COLM*, 2025. [URL: https://arxiv.org/abs/2504.13203](https://arxiv.org/abs/2504.13203).
- 6 A. Suvarna, C. A. Chance, K. Naranjo, *et al.*, “ModelCitizens: Representing community voices in online safety,” in *EMNLP*, 2025. [DOI: 10.18653/v1/2025.emnlp-main.1571](https://arxiv.org/abs/2504.15711).
- 7 A. Yerukola, S. Gabriel, N. Peng, and M. Sap, “Mind the gesture: Evaluating AI sensitivity to culturally offensive non-verbal gestures,” in *ACL*, 2025. [URL: https://arxiv.org/abs/2502.17710](https://arxiv.org/abs/2502.17710).
- 8 S. Gabriel, L. Lyu, J. Siderius, M. Ghassemi, J. Andreas, and A. Ozdaglar, “Generative AI in the Era of ‘Alternative Facts’,” in *Empirical Methods in Natural Language Processing (EMNLP)*, <https://arxiv.org/abs/2410.09949>, Nov. 2024.
- 9 S. Gabriel, I. Puri, X. Xu, M. Malgaroli, and M. Ghassemi, “Can AI Relate: Testing Large Language Model Response for Mental Health Support,” in *Empirical Methods in Natural Language Processing (EMNLP) Findings*, Nov. 2024. [URL: https://api.semanticscholar.org/CorpusID:269921604](https://api.semanticscholar.org/CorpusID:269921604).

Patents

- 1 H. Palangi, S. Gabriel, T. Hartvigsen, D. Ray, M. Sap, and E. Kamar, *Adversarial language imitation with constrained exemplars*.

PhD Committees

UCLA Advisees	■ Ashima Suvarna - Sheriff Issaka - Genglin Liu - Salman Rahman - Mehrab Beikzadeh - Elizabeth Eyeson
UCLA CS	■ Elaine Yixin Wan (Winter 2026) - Haoyi Qiu (Fall 2025) - Wenhan Yang (Spring 2025) - Christina Chance (Spring 2025) - Yufei Tian (Spring 2025) - Di Wu (Fall 2024)
UCLA ECE	■ Natarajan Balaji Shankar (Fall 2024)

Teaching

Instructor	■ Graduate Natural Language Processing (CS 263), 102 students. University of California Los Angeles, Winter 2026. - Inside the Black Box: Artificial Intelligence Safety and Mechanistic Interpretability, 34 students. University of California Los Angeles, Fall 2025. - Undergraduate Natural Language Processing (CS 162), 110 students. University of California Los Angeles, Spring 2025-26. - Computational Ethics, Large Language Models and the Future of NLP (CS 269), 47 students. University of California Los Angeles, Winter 2025. - Data Science Capstone (DS-GA 1006), 160 students. New York University, Fall 2023.
------------	--

Teaching (continued)

Guest Lecturer	- Black Studies in STEM. Caltech, Spring 2026. - ML for Healthcare. UC Berkeley, Spring 2025. - Speech Processing. University of California Los Angeles, Fall 2024. - Ethical Machine Learning In Human Deployments. Massachusetts Institute of Technology, Spring 2024. - Computational Ethics. Carnegie Mellon University, Spring 2023, 2024 and 2025. - Natural Language Processing. Massachusetts Institute of Technology, Fall 2023 and 2024. - AI Ethics. Oakton College, Fall 2023. - Natural Language Processing (Undergrad). University of Washington, Spring 2023. - Intro to Machine Learning. University of Washington, Fall 2022. - Natural Language Processing. Mount Holyoke College, Fall 2019. - NLP State-of-the-Art Methods. Carlson School of Management, Fall 2019.
Teaching Assistant	- Natural Language Processing (Undergrad/Grad). University of Washington. - Real Analysis. Mount Holyoke College.

Selected Recent Talks & Panels

- Invited talk at UW NLP seminar.
- Invited talk at UToronto SRI seminar.
- Panelist at LA Tech Week 2025 (Tech St Santa Monica).
- Invited talks at COLM 2025 workshops on Social Simulations with LLMs (keynote) and Visions of Language Modeling.
- Invited talk at the Inaugural Conference of the International Association for Safe and Ethical AI, February 2025
- Invited talks at Stanford NLP seminar, December 2021 and February 2025.
- Invited talk at UCLA Statistics and Data Science seminar, March 2024.
- Invited talk at NYU Center for Data Science Lunch seminar, November 2023.
- Invited talk at NYU-KAIST Inclusive AI Workshop, November 2023.
- Panelist for CHIL 2023 session on *LLMs and Healthcare*, June 2023.

Service

Area Chairing	ARR, COLM, ICLR, ICML, NeurIPS
Reviewing	ACL*, NAACL*, EMNLP, NeurIPS, AAAI, ICLR, ICML, Computational Linguistics, Journal of Artificial Intelligence, Journal of the American Medical Informatics Association, npj Digital Medicine, Nature Machine Intelligence (*Outstanding Reviewer)

Service (continued)

- | | |
|----------------------|--|
| Organizing Committee | <ul style="list-style-type: none">Faculty Advisory Committee for the Bedari Kindness Institute (2024-2026)Faculty Advisory Committee for the Ralph J. Bunche Center for African American Studies (2024-2027)UCLA CS MS Admissions Committee (2024-2026)Session Chair for Computational Social Science and Cultural Analytics (EMNLP 2024)NYU Data Science MS Admissions Committee (2023-2024)Tutorial Co-Chair (NeurIPS 2023)Socio-Cultural Inclusion Co-Chair and Generation Session Chair (NAACL 2022)Safety for E2E Conversational AI Special Session (SIGDIAL 2021) |
| Peer Mentoring | <ul style="list-style-type: none">GEM Computer Science Mentor at Mount Holyoke College (Spring 2016) |